

Lab_visualisation_questions

October 3, 2019

1 Data visualisation in Python

Prof. Dorien Herremans

For a full overview of types of plots using matplotlib, see the gallery at <https://matplotlib.org/2.0.2/gallery.html>

1.1 Scatterplots

We will be using `scottish_hills.csv` from <https://github.com/ourcodingclub/CC-python-pandas-matplotlib>. The file contains all the mountains above 3000 feet (about 914 metres) in Scotland.

We can read this into a variable and see the first 10 lines:

```
In [9]: import pandas as pd
```

```
url = "https://raw.githubusercontent.com/ourcodingclub/" \
      "CC-python-pandas-matplotlib/master/scottish_hills.csv"
dataframe = pd.read_csv(url)
print(dataframe.head(10))
```

	Hill Name	Height	Latitude	Longitude	Osgrid
0	A' Bhuidheanach Bheag	936.0	56.870342	-4.199001	NN660775
1	A' Chailleach	997.0	57.693800	-5.128715	NH136714
2	A' Chailleach	929.2	57.109564	-4.179285	NH681041
3	A' Chraileag (A' Chralaig)	1120.0	57.184186	-5.154837	NH094147
4	A' Ghlas-bheinn	918.0	57.255090	-5.303687	NH008231
5	A' Mhaighdean	967.0	57.719644	-5.346720	NH007749
6	A' Mharconaich	973.2	56.857002	-4.290668	NN604762
7	Am Basteir	934.0	57.247931	-6.202982	NG465253
8	Am Bodach	1031.8	56.741727	-4.983393	NN176650
9	Am Faochagach	953.0	57.771801	-4.853899	NH303793

As explored last week, pandas dataframes can be used for some preliminary data exploration. For instance, let's sort the hills by height:

```
In [10]: sorted_hills = dataframe.sort_values(by=['Height'], ascending=False)
```

```
# Let's have a look at the top 5 to check
print(sorted_hills.head(5))
```

	Hill Name	Height	Latitude	Longitude	Osgrid
92	Ben Nevis	1344.5	56.796891	-5.003675	NN166712
88	Ben Macdui (Beinn Macduibh)	1309.0	57.070368	-3.669099	NN988989
104	Braeriach	1296.0	57.078298	-3.728389	NN953999
115	Cairn Toul	1291.0	57.054397	-3.710773	NN963972
212	Sgor an Lochain Uaine	1258.0	57.058369	-3.725797	NN954976

Now let's load matplotlib. Note: if you are using a jupyter notebook you need the inline statement on line 1 below:


```
In [0]: %matplotlib inline
import matplotlib.pyplot as plt
```

To save us some time, let's create some shortcut variables, x and y, to register the hight and latitude coordinates of each of the hills.

```
In [0]: x = dataframe.Height
        y = dataframe.Latitude
        z = dataframe.Longitude
```

Now we are ready to start visualising them. Let's create (and save) a scatterplot:

```
In [13]: plt.scatter(x, y)
plt.savefig("scottish_scatter_plot.png")
```



If you are not using iPython, you can use `plt.show()` to display the plot.
Now let's build upon this graph by adding a linear regression line to it.


```
In [0]: from scipy.stats import linregress
stats = linregress(x, y)

m = stats.slope
b = stats.intercept
```

Now we can add the plot of our linear regression by using the equation of a straight line:

```
In [15]: plt.scatter(x, y)
# The equation of the straight line.
plt.plot(x, m * x + b, color="red")
```

```
Out[15]: [<matplotlib.lines.Line2D at 0x7f87b9e04668>]
```



Note, whether this line is statistically significant can be determined using the extra information in the stats object - `stats.rvalue` and `stats.pvalue`.

Now you can make your plot look nicer using arguments such as `fontsize`, `linewidth`, `color`,...

```
In [16]: # Change the default figure size
plt.figure(figsize=(10,10))
```

```

# Change the default marker for the scatter from circles to x's
plt.scatter(x, y, marker='x')

# Set the linewidth on the regression line to 3px
plt.plot(x, m * x + b, color="red", linewidth=3)

# Add x and y labels, and set their font size
plt.xlabel("Height (m)", fontsize=20)
plt.ylabel("Latitude", fontsize=20)


# Set the font size of the number labels on the axes
plt.xticks(fontsize=18)
plt.yticks(fontsize=18)

```

```

Out[16]: (array([56. , 56.5, 57. , 57.5, 58. , 58.5, 59. ]),
<a list of 7 Text yticklabel objects>)

```




Let's have a look at how the hills are spread out geographically using latitude (y) and longitude (z). Now we can use $s = x$ to say that the size needs to be equal to the height (x). (I added -900 to make the difference between big and small mountains larger)

```
In [17]: import numpy as np
```

```
colors = np.random.rand(len(y)) #generates a different color  
                                         # for each different mountain
```

```
plt.scatter(y, z, s = (x-900), c=colors, alpha=0.5)
```

```
Out [17]: <matplotlib.collections.PathCollection at 0x7f87b877a320>
```




1.2 Histograms

Let's try some other graphs. A full selection is given at the [matplotlib website](https://matplotlib.org/).

Let's start by looking at the distribution of our hills over the latitude (variable y from earlier):

```
In [19]: plt.hist(x, bins=25, density=True) #bins separates the latitude in  
                                         # 25 discrete categories. Density will normalize the data to 1.  
                                         # If you get an error, use density instead of normed (newer matplotlib version)  
  
plt.savefig("histogram.png", dpi=25) # results in 160x120 px image
```



Quickly style your plot with stylesheets, full overview at https://matplotlib.org/gallery/style_sheets/style_sheets_reference.html.

Let's also create a new variable that contains the height of the hills -100m. This to illustrate how to add a second distribution to your graph. In this case, we will make them slightly transparent.

In [20]: `import numpy as np`


```
# using a stylesheet:
plt.style.use('seaborn-pastel')

#creating the new height variable:
shifted_x = x - 100

fig, ax = plt.subplots()
ax.hist(x, bins=25, density=True, histtype="stepfilled",
        alpha=0.8, label='Height')
ax.hist(shifted_x, bins=25, density=True, histtype="stepfilled",
        alpha=0.8, label='Height - 100')
ax.legend(prop={'size': 10})

ax.set_ylabel('Normalised distribution')
ax.set_xlabel('Height')
```

Out[20]: `Text(0.5, 0, 'Height')`



1.3 Bar charts

Plot the average CO2 output of both vegetarians and meat eaters for different continents in a bar plot. The data is given below:

```
In [21]: labels = ['EU', 'US', 'AS']
```

```
meateaters = [122, 135, 80]
```

```
vegetarians = [40, 43, 23]
```

```
# Keep a numeric index for the x-axis labels. This will be the position of the ticks
index = np.arange(len(labels))
```

```
fig, ax = plt.subplots()
```

```
bar_width = 0.35
```

```
opacity = 0.6
```

```
rects1 = ax.bar(index, meateaters, bar_width,
                 alpha=opacity, color='b',
                 label='Meat eaters')
```

```
# Note, shift the x position also with the bar width so that the bar
```

```

# appears next to the previous ones
rects2 = ax.bar(index + bar_width, vegetarians, bar_width,
                alpha=opacity, color='r',
                label='Vegetarians')

ax.set_xlabel('Continent')
ax.set_ylabel('CO2 emissions')
ax.set_title('CO2 emissions per continent, per diet')

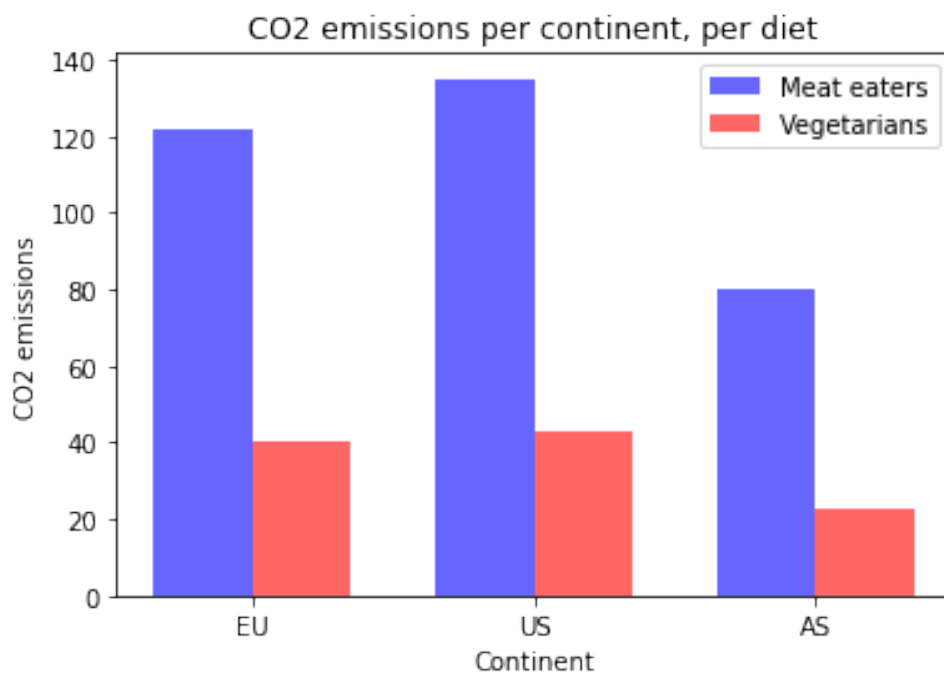
ax.set_xticks(index + bar_width / 2)

ax.set_xticklabels(labels)

ax.legend()

```

Out[21]: <matplotlib.legend.Legend at 0x7f87b8545160>



1.4 Line plots

Let's move on to another type of graph: a simple line plot, but using two vertical axis. We will create a function to calculate the temperature in celcius given Fahrenheit.

```

In [0]: def fahrenheit2celsius(temp):
        """

```



```

Returns temperature in Celsius.
"""
return (5. / 9.) * (temp - 32)

# TEMPURATURE AT EACH HOUR THROUGHOUT THE DAY (In Fahrenheit)
temperature = [100, 102, 106, 105, 90, 85, 85, 89, 100, 102, 103, 108, 100, 102,
               106, 105, 90, 85, 85, 89, 100, 102, 103, 108]

```

We will use `twinx` to get a second set of axes. This allows us to plot the temperature evolution throughout the day using both C and F.

```

In [23]: fig, ax_f = plt.subplots()
         ax_c = ax_f.twinx()

         #plot our data:
         ax_f.plot(temperature)
         ax_f.set_xlim(0, 24) #x-axis shows 24 hours

         # set the axis limits:
         y1, y2 = ax_f.get_ylim()
         ax_c.set_ylim(fahrenheit2celsius(y1), fahrenheit2celsius(y2))
         ax_c.figure.canvas.draw()


         # change some axis labels
         ax_f.set_title('Two scales: Fahrenheit and Celsius')
         ax_f.set_ylabel('Fahrenheit')
         ax_c.set_ylabel('Celsius')
         ax_f.set_xlabel('Time (hour of day)')

```

```

Out[23]: Text(0.5, 15.0, 'Time (hour of day)')

```




Suppose we have a sinusoidal plot (formula line 2 below), and we want to fill it. First let's plot the sinusoid.

```
In [24]: newx = np.linspace(0, 1, 500) # sample 500 X's between 0 and 1
        newy = np.sin(4 * np.pi * newx) * np.exp(-5 * newx) #formula for our graph

        fig, ax = plt.subplots() #plt.subplots() lets us acces the axis
        # and plot seperately.

        ax.plot(newx, newy)
```


```
Out[24]: [<matplotlib.lines.Line2D at 0x7f87b8501b70>]
```



Then we can easily use the fill command.

```
In [25]: fig, ax = plt.subplots()
         ax.fill(newx, newy)
```

```
Out[25]: [matplotlib.patches.Polygon at 0x7f87b856c4e0>]
```




Axis labels can influence how we perceive the data. Let's have a look at this stock, which has been stagnating recently.

```
In [26]: # growth of stock A
stockA = [0, 10, 20, 30, 40, 60, 80, 140]

plt.plot(stockA)
# plt.ylim(bottom=0, top = 110)
plt.xlabel('Time')
plt.ylabel('Stock A price')
```

```
Out[26]: Text(0, 0.5, 'Stock A price')
```



Now let's change to a logarithmic axis. Other options here are linear, log, logit, symlog. Give it a try.


```
fig = plt.figure() ax = fig.add_subplot(1, 1, 1)
line, = ax.plot(stockA) ax.set_yscale('log') #change the scale here
plt.xlabel('Time') plt.ylabel('Stock A price') ax.set_title('Logarithmic')
This illustrates how axes can deform the data...
```

1.4.1 Boxplots

For all the mountains, let's see what their average height is, with standard deviation in a boxplot.

```
In [27]: plt.boxplot(x)
```

```
Out[27]: {'boxes': [<matplotlib.lines.Line2D at 0x7f87b82f76d8>],  
          'caps': [<matplotlib.lines.Line2D at 0x7f87b82f7ef0>,  
                  <matplotlib.lines.Line2D at 0x7f87b8301278>],  
          'fliers': [<matplotlib.lines.Line2D at 0x7f87b8301908>],  
          'means': [],  
          'medians': [<matplotlib.lines.Line2D at 0x7f87b83015c0>],  
          'whiskers': [<matplotlib.lines.Line2D at 0x7f87b82f7828>,  
                      <matplotlib.lines.Line2D at 0x7f87b82f7ba8>]}
```




We can make this slightly nicer:



```
In [28]: # notched plot  
plt.figure()  
plt.boxplot(x, 1)  
  
# change outlier point symbols  
plt.figure()  
plt.boxplot(x, 0, 'gD')  
  
# don't show outlier points  
plt.figure()  
plt.boxplot(x, 0, '')  
  
# horizontal boxes  
plt.figure()
```



```
plt.boxplot(x, 0, 'rs', 0)

# change whisker length
plt.figure()
plt.boxplot(x, 0, 'rs', 0, 0.75)
```

```
Out[28]: {'boxes': [<matplotlib.lines.Line2D at 0x7f87b81bf9b0>],
'caps': [<matplotlib.lines.Line2D at 0x7f87b81cc208>,
<matplotlib.lines.Line2D at 0x7f87b81cc550>],
'fliers': [<matplotlib.lines.Line2D at 0x7f87b81ccbe0>],
'means': [],
'medians': [<matplotlib.lines.Line2D at 0x7f87b81cc898>],
'whiskers': [<matplotlib.lines.Line2D at 0x7f87b81bfb00>,
<matplotlib.lines.Line2D at 0x7f87b81bfe80>]}
```







2 Now to try yourself:

Load the dataset from https://raw.githubusercontent.com/plotly/datasets/master/school_earnings.csv, and have a look at what it contains.

Then create the following:

1. A **histogram** of the salaries for women.
2. Add the men's salaries to this histogram.
3. Give your histogram a dark background and label the axes.
4. Next, please label the colors of the histogram so we know who is what (men vs women).
5. Instead of a histogram, create a **bar chart** that lists the salary for women (y-axis) for each school.
6. Also add men to this bar chart.
7. Make the style nice and add labels.
8. Now create a nice **boxplot** of the data, one for men, one for women (two box's same graph).

2.1 Solution: histogram

In [0]:

In [0]:

2.2 Solution: Bar chart

In [0]:

2.3 Solution: Boxplot

In [0]:

Well done!

In [0]: